ZAVA FORWARD

More homes ready. Faster. With control intact.

One workbench for the retrofit programme, with damp, mould and hazard referral built in

Decision today: scope, owners, then a costed plan



Case Study 48 · Zava North Homes is fictional

ACT 1 · THE CHALLENGE

15,600 homes to lift by 2030. Thousands more to keep safe.

15,600
homes below EPC C
26% of social homes

3,600
homes with damp
6% of housing association homes

22,800
households likely to report condensation, damp or mould
38%

3,210
stage-one complaints a year
53.5 per 1,000 homes

Sources: English Housing Survey 2024-25 · RSH Tenant Satisfaction Measures 2024/25 · national rates applied to 60,000 homes



The dates matter. So does their legal status.

Dated source milestones, not five interchangeable legal duties.



Awaab's Law Phase 1

In-force regulations in retained register

Cold and heat Phase 2 guidance

Planned expansion; check commencement

Remaining hazards Phase 3 plan

Planned scope; not a live rule

EPC C Policy direction

Applicability and exceptions need review

Decent Homes Policy direction

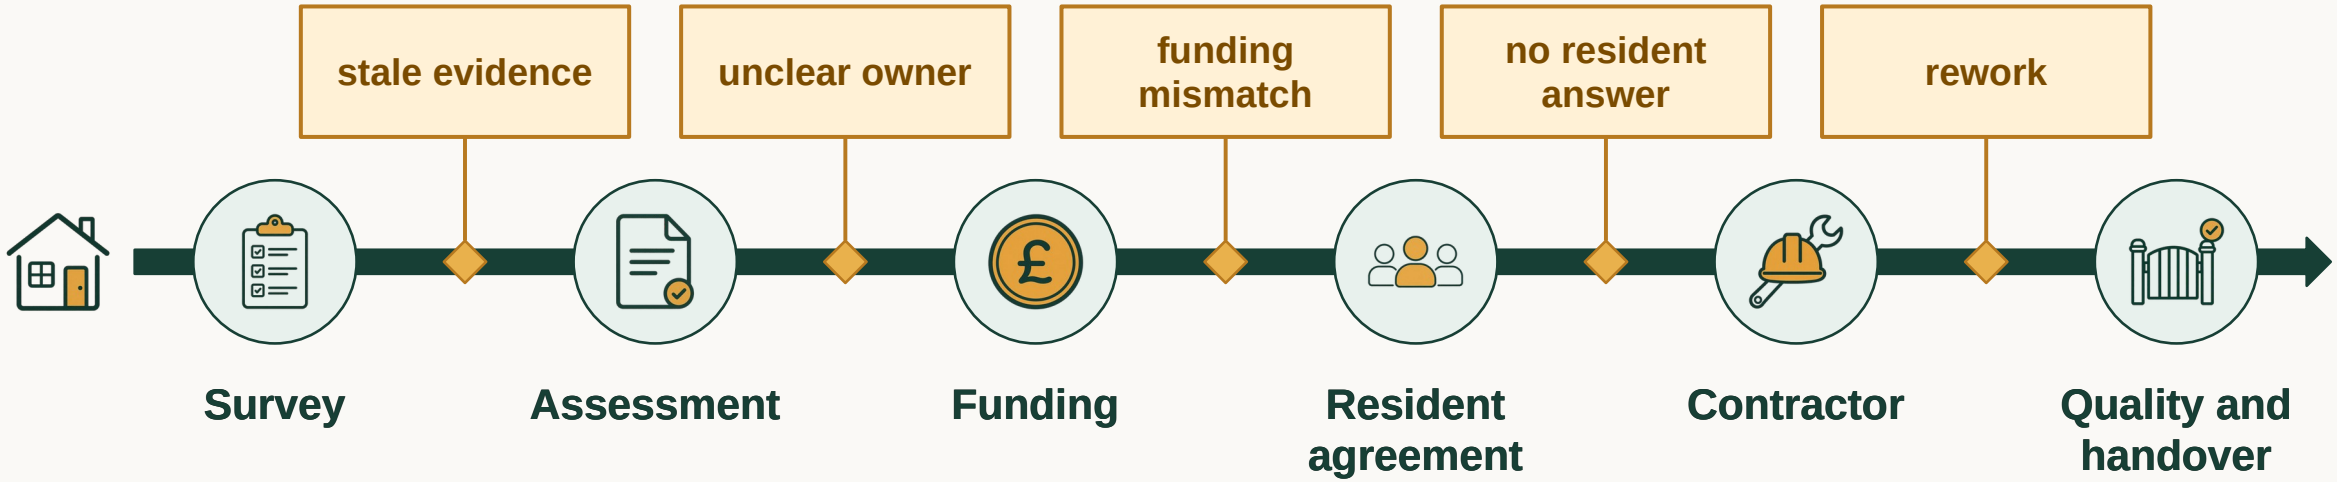
Implementation not proved here

One connected service. Distinct legal decisions.

Verify current official sources and competent review before activating a rule.

Sources: retained regulatory register, 19 Sep / updated 15 Oct 2026. No fresh legal-currentness check.

Homes don't stall on ambition. They stall between teams.



Evidence, funding, capacity and resident choice must align

More starts are not the same as more accepted homes

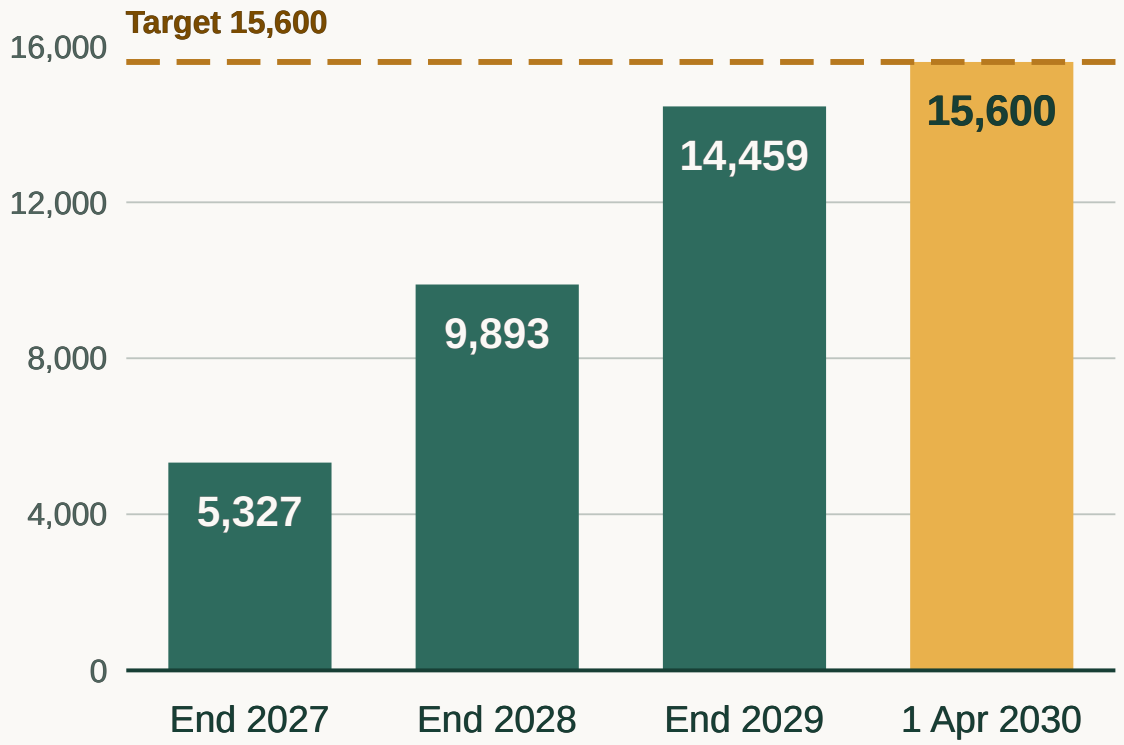
A conflicting record should stop one approval, not all useful work

Sources: exercise journeys, role research and agent-request challenge

Forecast: every Zava home at EPC C by April 2030.

FORECAST

Homes at EPC C, cumulative from Nov 2026



15,600 homes · 41 months

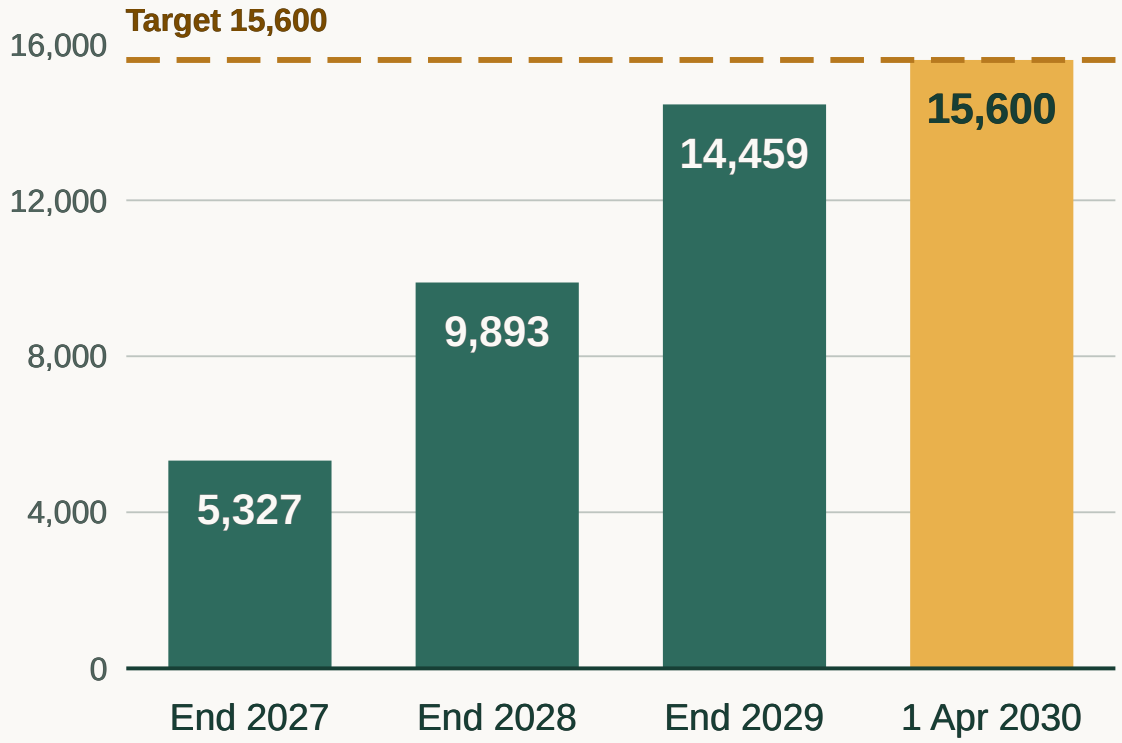
88 homes a week

Forecast · EHS 2024-25 · SRS MEES impact assessment · Warm Homes: Social Housing Fund wave 3 · assumptions in notes

Forecast: every Zava home at EPC C by April 2030.

FORECAST

Homes at EPC C, cumulative from Nov 2026



15,600 homes · 41 months

88 homes a week

GBP 95m
works benchmark

GBP 47.4m
conditional grant potential

USD 6,873
whole-cohort inference scenario

Forecast: 15,600 homes; E21: assumed workload + secondary price (23 Oct 2026). Grant subject to eligibility, allocation, award and claim.



ACT 2 · DISCOVERY

Discovery grew scope. Delivery keeps learning.

Activity and scope, not outcomes

1147 non-merge commits	**982** merged pull requests	**260** tracked project items	**126** completed issues
125 open issues	**9** not planned	**28** journey hypotheses	**5** proposal capabilities

Snapshot: 8 Nov 2026

Ongoing build; activity is not customer benefit

Sources: git bcd86707; GitHub merged PRs; AMA Plan 7, 8 Nov 2026, 14:12 UTC. Refresh before final delivery.

ACT 2 · DISCOVERY

From the brief to bounded design.

01 Ground the brief → **02 Ask: what if?** → **03 Map the work**

Case facts + official context
Company unknowns stay unknown

8 role instruments: richer evidence, photos and contractor capacity

9 later instruments / 28 maps
Needs, pains, gains and handoffs

PAIN HYPOTHESIS -> DESIGN RESPONSE -> WHAT WE TEST

Repeated story

Owned case + optional access

Test: Effort and inclusion

False commitment

Offer != booking; choice stays

Test: Cancellation and recovery

Changed evidence

Version-bound review

Test: Only affected reviews reopen

Desk research, owner feedback and synthetic tests; direct resident validation remains pending.

Sources: original brief; sector evidence; role challenge; journey atlas; DT Coach decision trace (30 Oct).

ACT 2 · DISCOVERY

We chose a controlled workflow with bounded AI.

BRIEF SAID

Five tasks with explicit L2-L4 levels

WE BUILT

Five L2 evidence proposals; fuller scope remains

Choice: rules + bounded AI. Rules-only stays the comparator.

BRIEF SAID

Programme supervision and coordination

WE BUILT

Shared case, workflow and accountable handoffs

Why: shared state first; live connectors remain gated.

BRIEF SAID

People approve; residents control consent

WE BUILT

The same decision rights, enforced in the API

This meets the brief's boundary; it is not a new restriction.

BRIEF SAID

Resident channel with consent management

ADDED OPTION

Optional website + advice; assisted routes stay

Why: easier own-case answers, never digital-only service.

Sources: agent-request challenge (E03) · ADR 0002 and ADR 0003 (E06, E08, E09)

S08 · Answers: B03 · B14

Every prototype taught us something we kept.

26 Sep
PR #1 requirements review
open decisions made explicit

7 Oct
Owner: clearer case context
next owner and plain status

8 Oct
Prototype simplicity review rules first; AI adds drafting
total review effort unknown
retain the simpler comparator

14 Oct
28 journey hypotheses
131 tracked issues created

14 Oct
Five real proposals, 6,884 tokens + SQL restart
integration proved;
human value still unknown

27 Oct
Optional resident site, own-case help and synthetic sensors
assisted routes remain

8 Nov
Source and release hardening
125 tracked issues still open; final evidence refresh ahead



Brief -> options -> prototypes -> backlog -> build -> review

Scope growth records discovery; closed issues are not housing outcomes

Next: accepted journeys, useful assistance and tested autonomy

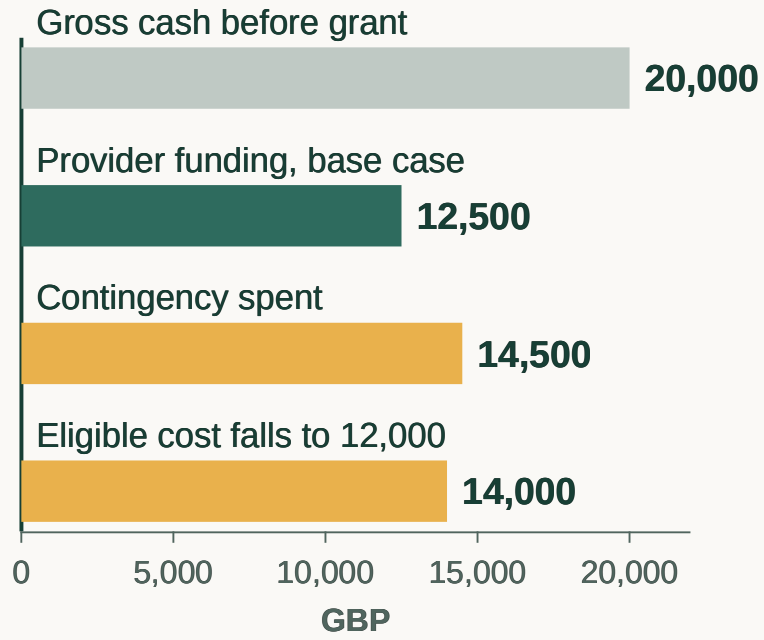
Sources: PRD revision history and learnings · exercise addendum · prototype learnings review · pull requests #1, #19, #58, #433

WHAT-IF

We stress-tested the economics before we wrote the product.

20 scenarios
five agents, a programme supervisor and a damp and mould loop

ONE HOME · exercise fixture (fictional)



THE PROGRAMME · 15,600 homes

WORKS

GBP 6,082/home historical HA benchmark — **94.9m**

By EPC band — **£107m**

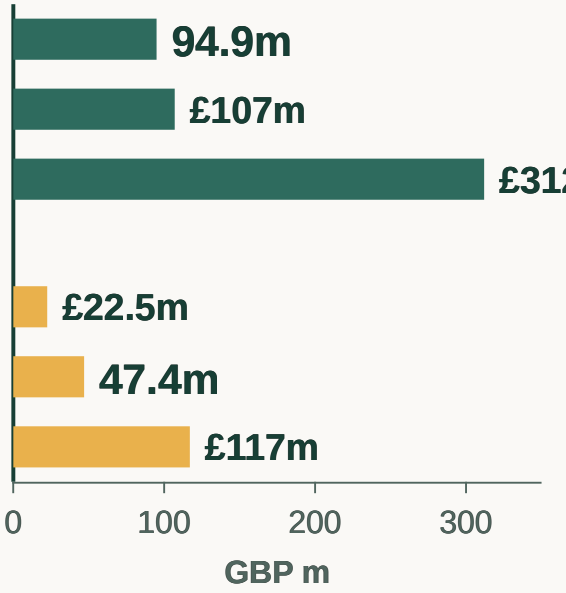
Deep-retrofit fixture (GBP 20,000 a home) — **£312m**

GRANT SCENARIOS

Proportional wave 3 share — **£22.5m**

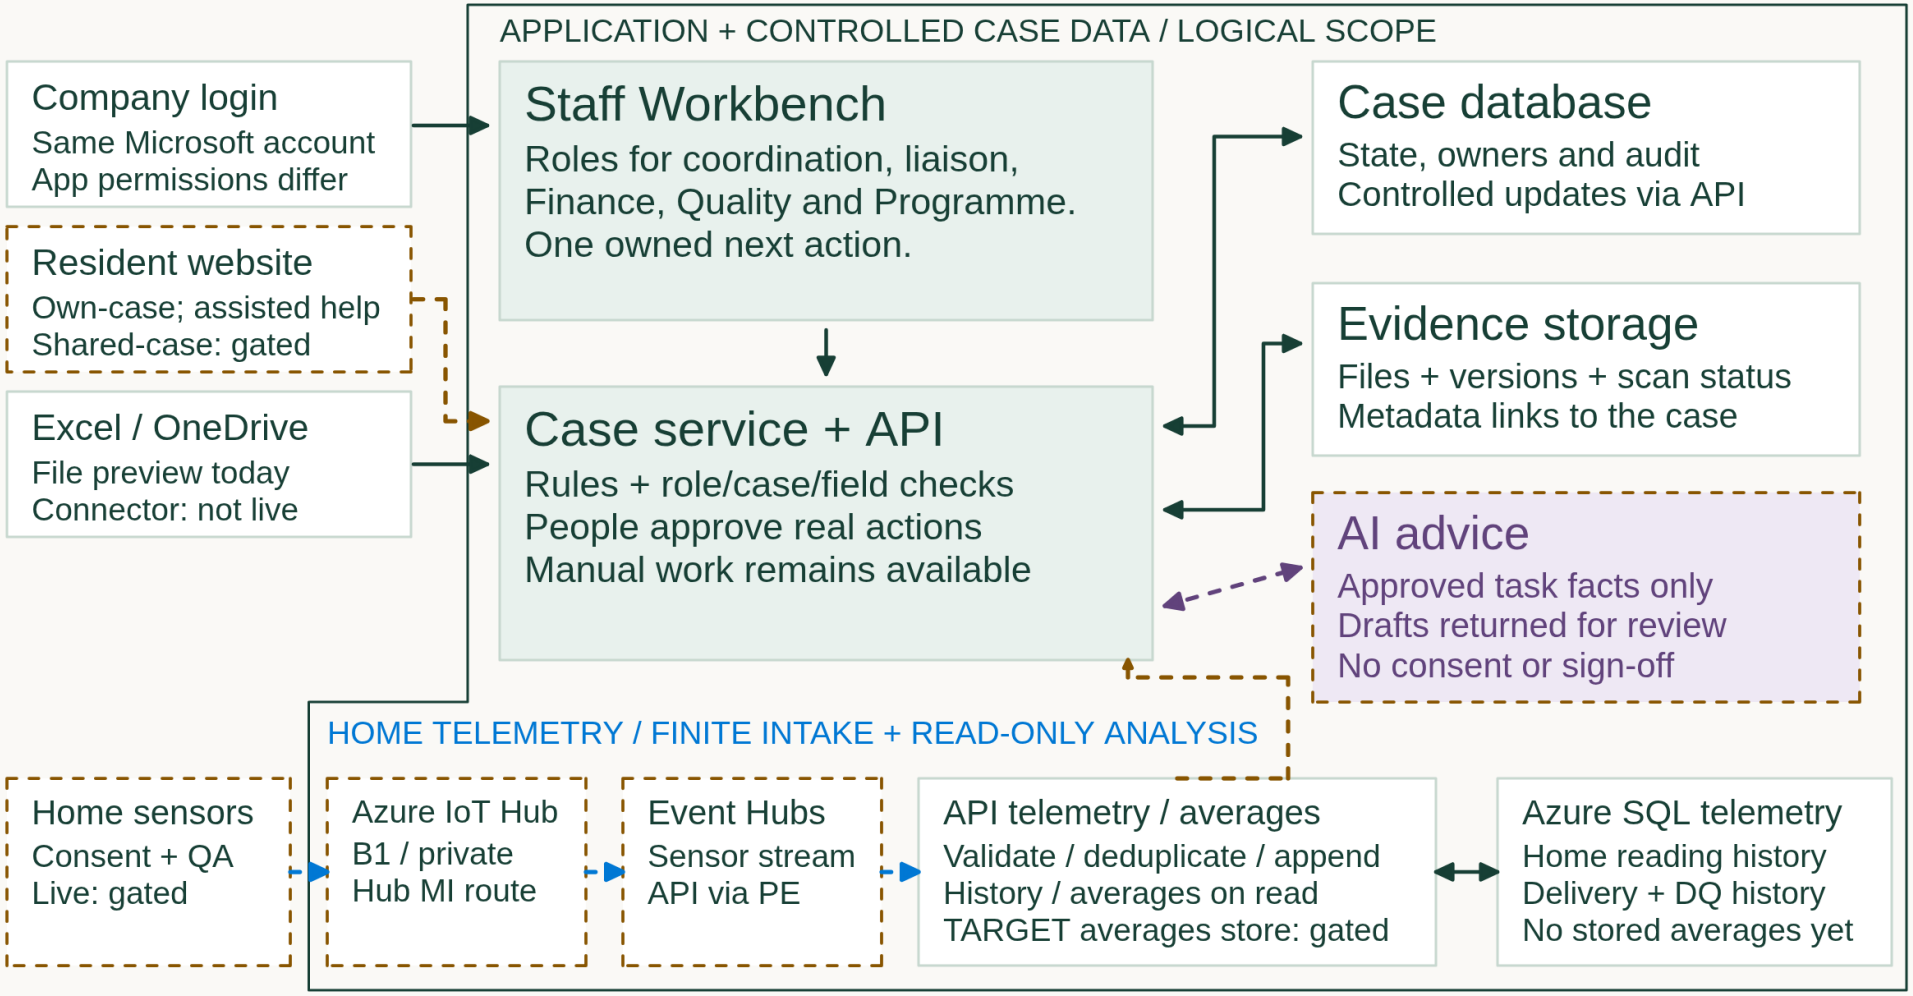
Conditional GBP 3,041/home (50% share) — **47.4m**

Full £7,500 cap — **£117m**



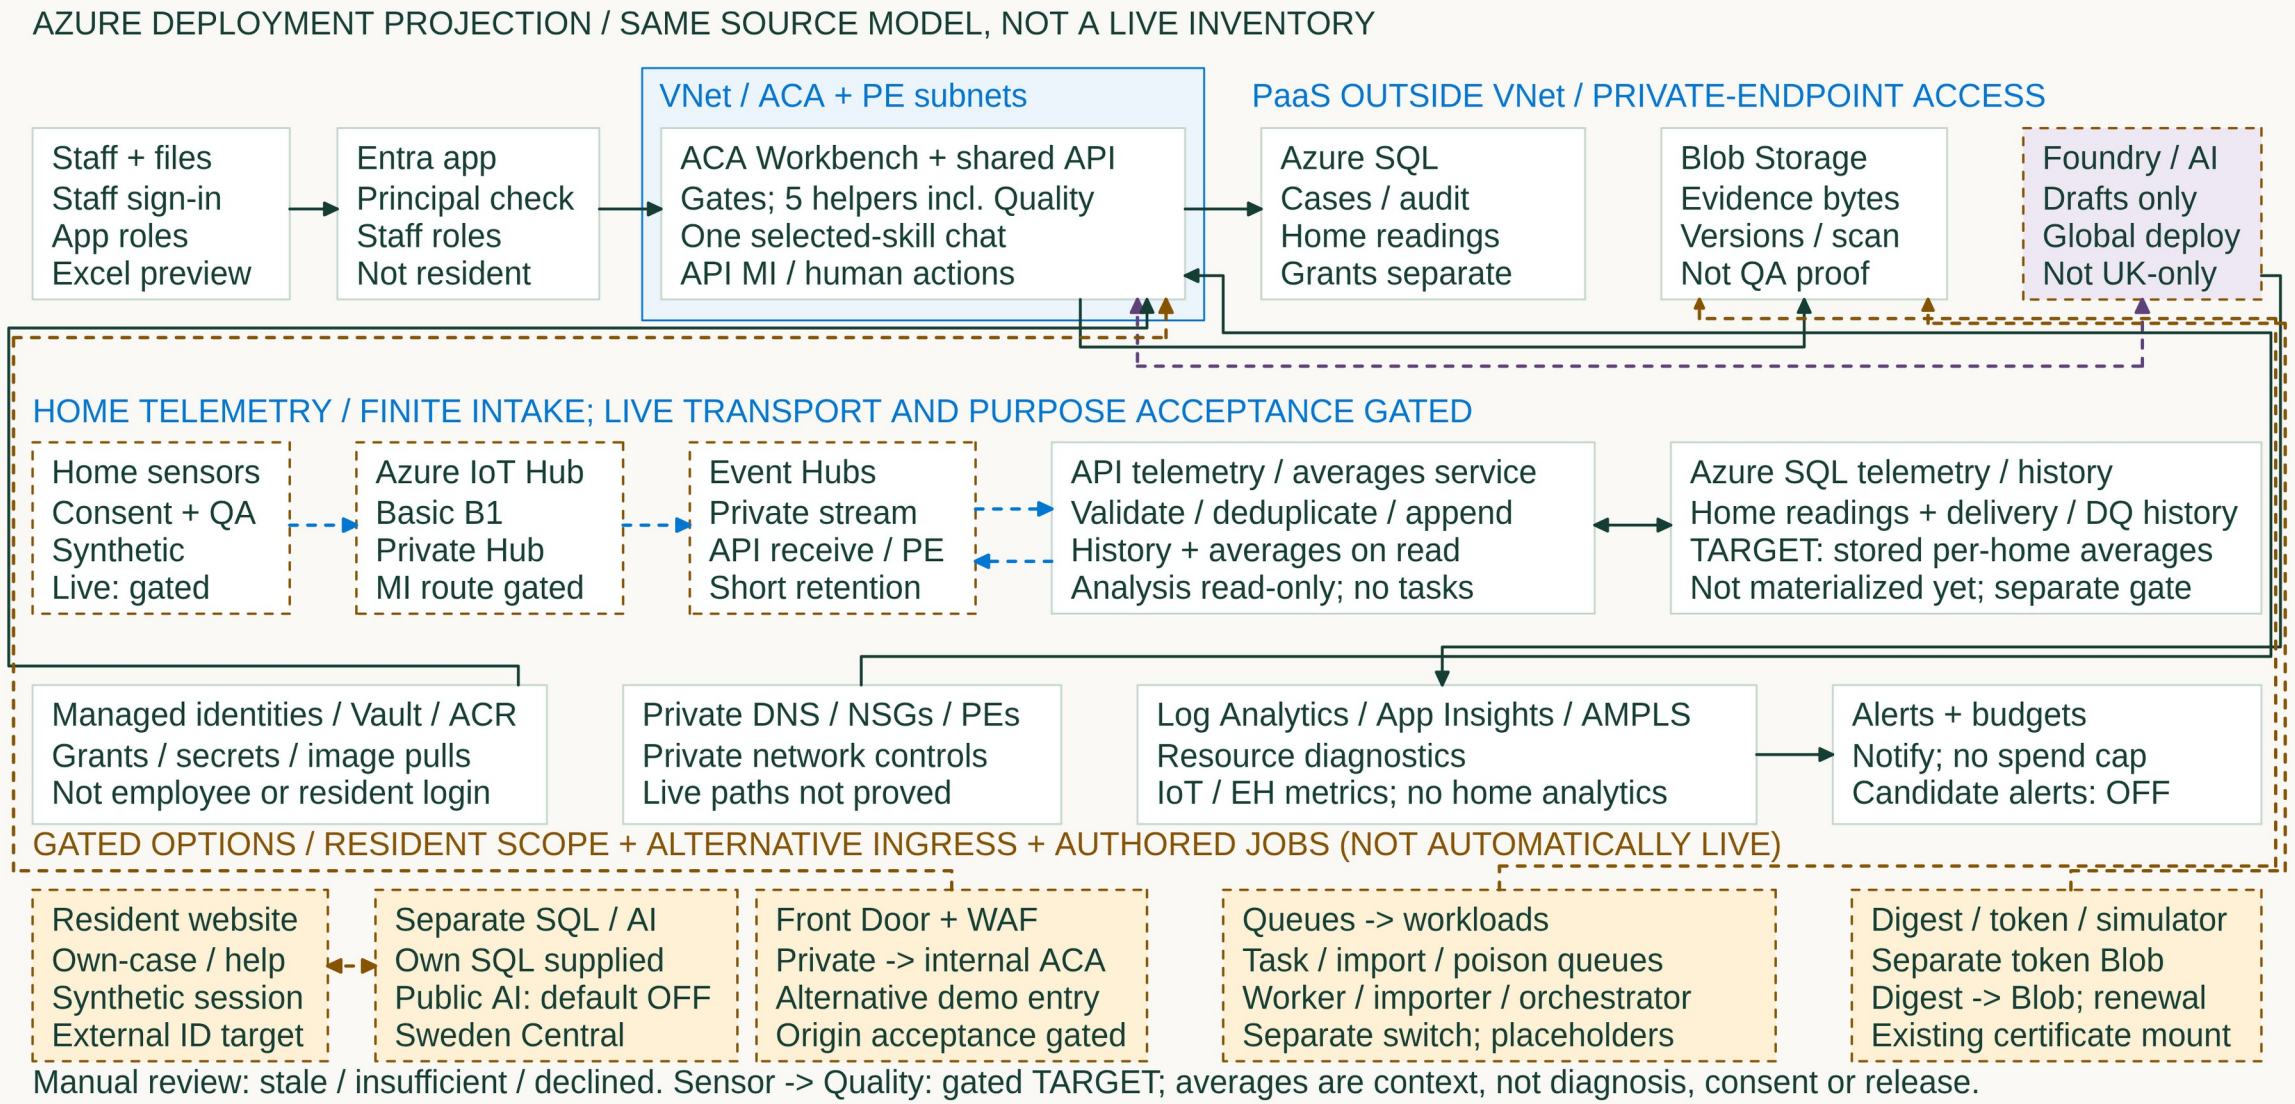
E21: dated benchmark, not a quote. Grant assumes eligible homes/costs, allocation, award and claim. GBP 20,000 fixture remains separate.

One case service connects people, evidence and advice.



Conceptual source-backed view, 8 Nov 2026. Dashed paths require separate acceptance; not a live inventory.

Azure connects home evidence to owned action.



A resident's journey: less friction, more choice.

Desk-research hypotheses; direct resident validation pending.

	01 Contact	02 Survey	03 Decide	04 Installation	05 Live with it
Human role	Housing	Assessor	Resident	Programme	QA / Housing
ACTION / NEED	Call / be heard	Plan visit	Choose works	Keep routines	Use / query
PAIN TO TEST	Story repeated	Needs missed	Rushed choice	Disrupted day	Concern lost
GAIN SOUGHT	Clear contact	Visit support	Pause / defer	Agreed visit	Owned help

Phone, letter and assisted visits stay. QA stays independent. Urgent safety goes straight to people.

Sources: resident journey atlas and DT Coach source-linked map. Editable Excalidraw + text equivalent retained.

Rules run the workflow. AI helps with the evidence.

WORKBENCH / ILLUSTRATIVE CASE

EXAMPLE-HOME-0412

Evidence conflict: two sources disagree

OWNED NEXT ACTION

Assessor: resolve the source conflict

Source-linked help, not an accepted case decision.

1 Rules choose the next task

Case facts -> coded gates -> next owner
No model is needed to keep work moving

2 AI prepares evidence

Staff request or approved scan -> proposal
Case card + queue; no case-state change

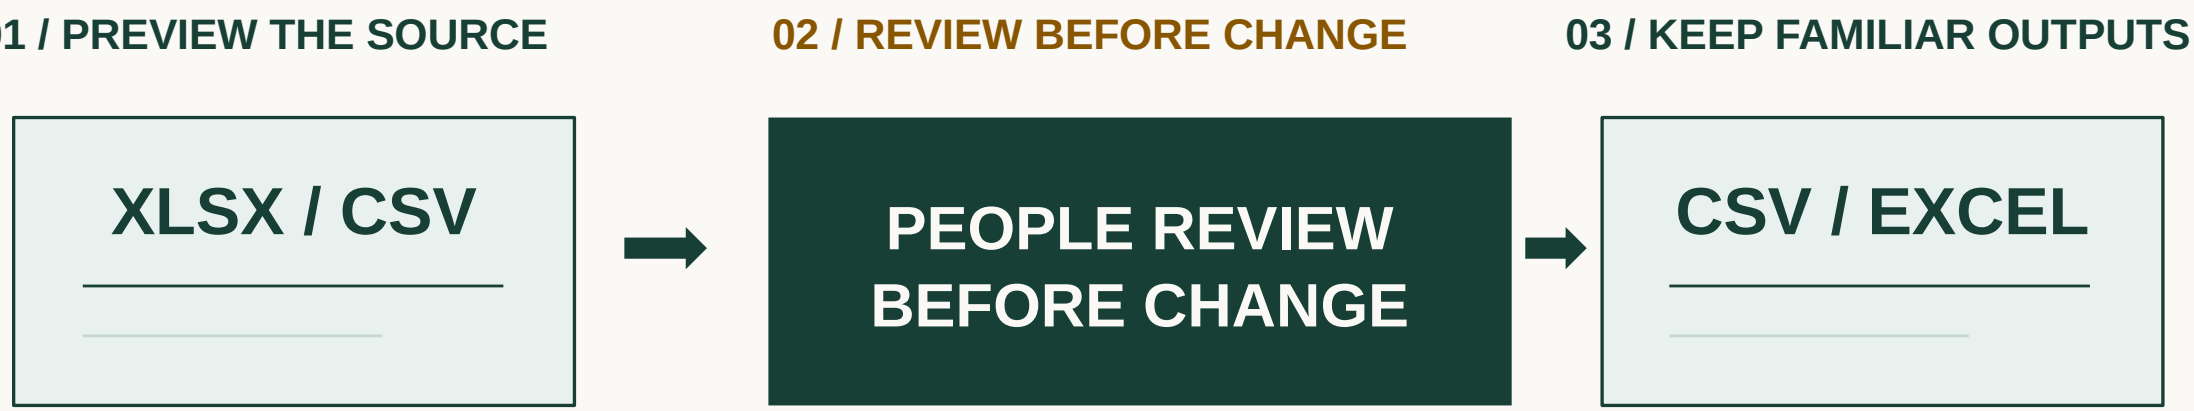
3 People review or ask

Check/edit/reject; optional follow-up chat
Consent and sign-offs stay separate

WORKFLOW IN CODE | AI is an optional branch, never the state-transition authority.

Sources: current API/core/agents contracts, 8 Nov 2026. Automatic preparation is not automatic approval.

Adopt without a big-bang replacement.



Declared workbook structure
Row and source lineage
Original file unchanged

Conflicts and quarantine
Separate human review
No guessed identity joins

Read-only comparisons
Familiar packs and exports
Manual and assisted routes

Existing file/CLI tooling. The web import journey is not accepted.

Owners accept mapping, reconciliation and rollout. A live enterprise connector is a separate decision.

Sources: importer/API guides. Existing CLI/file tooling, not an accepted web import journey.

Five responsibilities. Autonomy is not one switch.

S	A	F	C	Q
Stock	**Assessment**	**Funding**	**Works + residents**	**Quality**
L2 -> L2 recommend Rules + stock facts	L3 -> L2 preparation Survey/design checklist	L3 -> L2 preparation Application/claim gaps	L3 acts -> L2 options No auto-book or send	L4 flags -> L2 concerns No autonomous QV flag

RULES ENFORCE

Ranking, arithmetic,
explicit state-flow gates

DECISION 1 / HISTORICAL

Candidate / support scoring
Current rebuild still pending

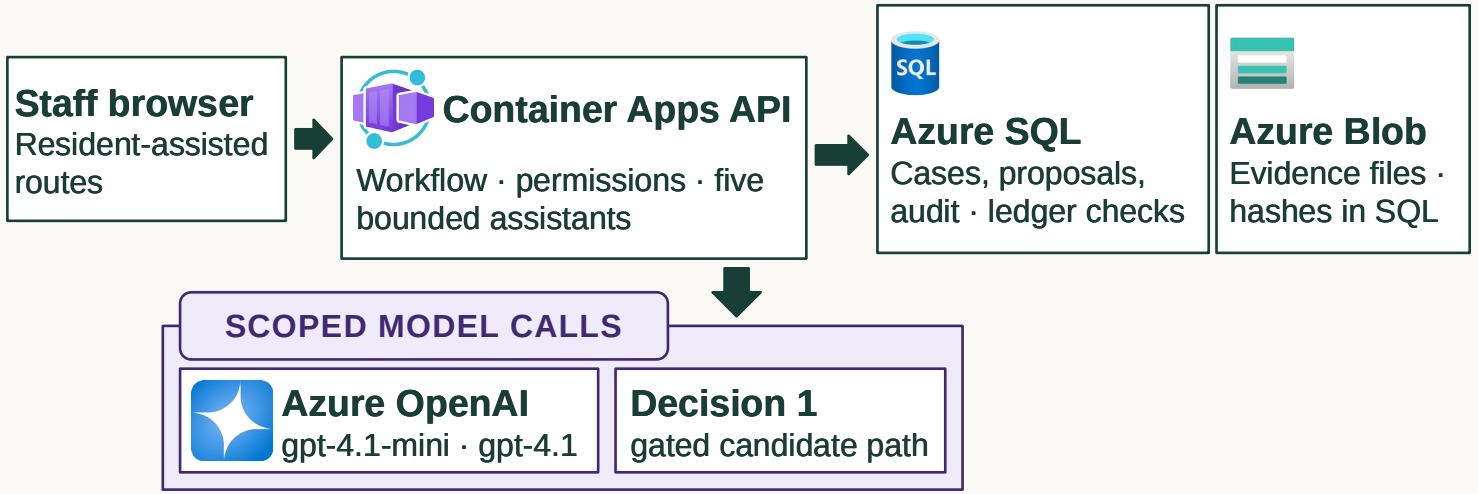
GENERATIVE AI DRAFTS

Five proposals + chat
Source + meaning review

PEOPLE DECIDE | Consent. Technical sign-off. Finance approval. Independent quality acceptance.

Brief -> current proposal authority. Full task scope and activation remain separate; B03 / page 30.

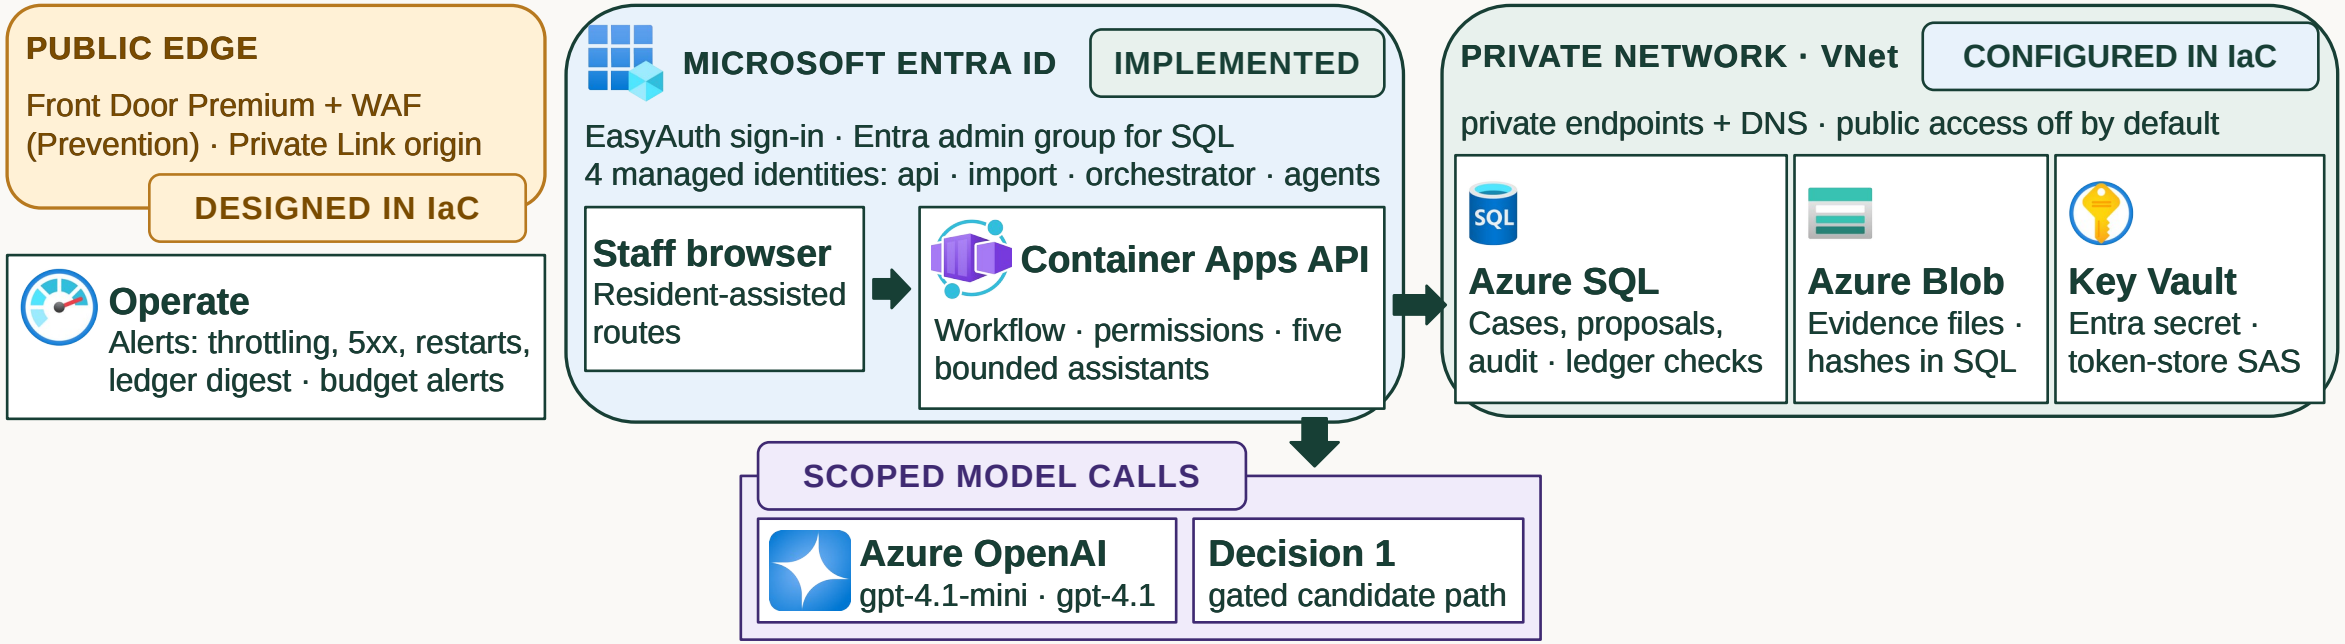
Secure by design: every home record sits inside three boundaries.



Sources: infra/main.bicep · architecture README · ADR 0002 decision 6 · PRD NFR-003 · BRD no-self-approval rules · AMA Plan #80, #102, #143

Filled arrow = request path · tinted region = security boundary · chip = delivery status

Secure by design: every home record sits inside three boundaries.



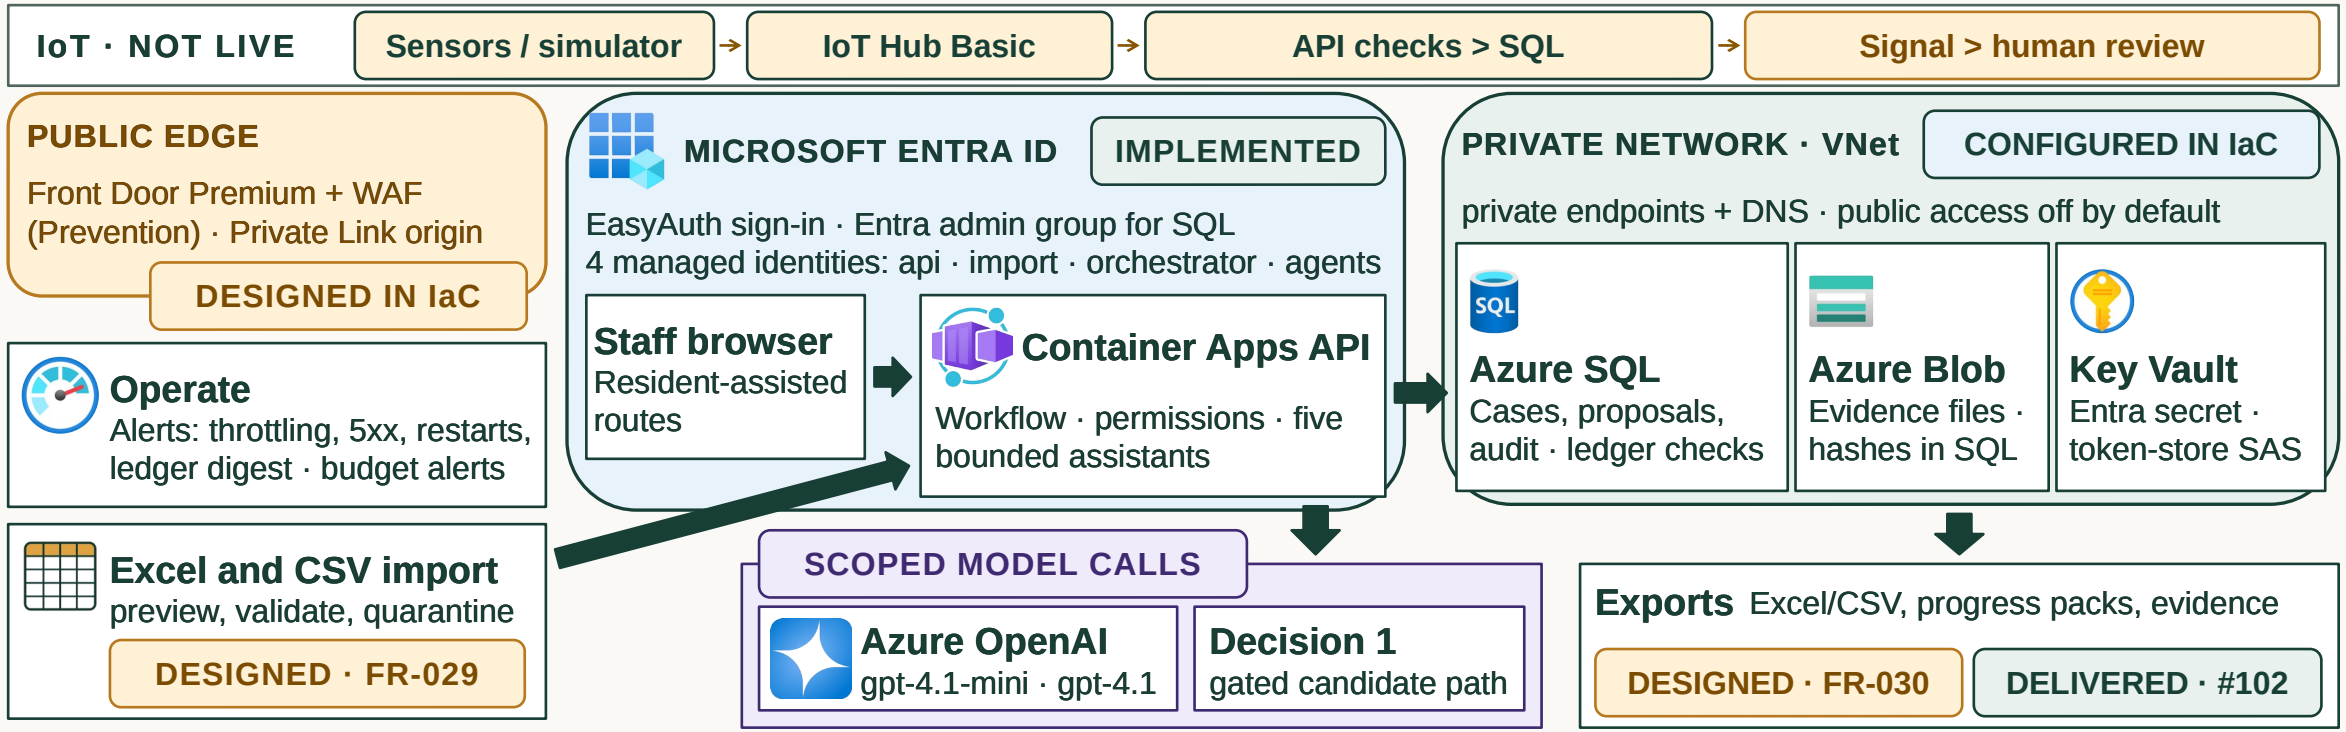
Current staff sign-in and Graph acceptance remain outstanding.

Network: SQL and Key Vault on private endpoints · managed identities · no SQL logins

Sources: infra/main.bicep · architecture README · ADR 0002 decision 6 · PRD NFR-003 · BRD no-self-approval rules · AMA Plan #80, #102, #143

Filled arrow = request path · tinted region = security boundary · chip = delivery status

Secure by design: every home record sits inside three boundaries.

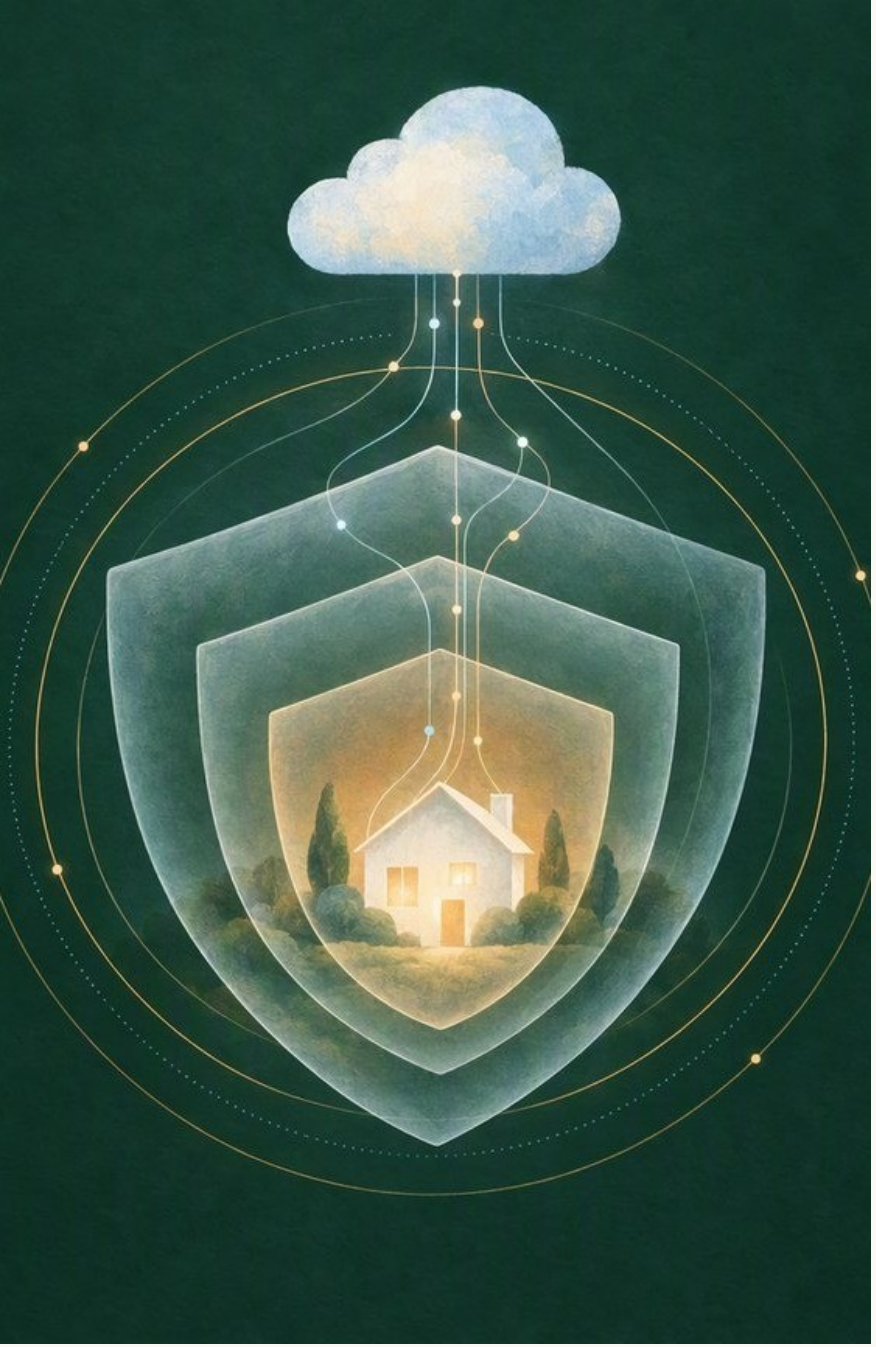


Current staff sign-in and Graph acceptance remain outstanding.

Network: SQL and Key Vault on private endpoints · managed identities · no SQL logins

IoT: monitor after retrofit; no diagnosis. Excel/M365 stay. Portal: assess #441.

Sources: ADR 0002/0003 · infra/main.bicep · PRD NFR-003 · #441 (assessment) · AMA Plan #80, #102, #143



ACT 4 · BUILT TO BE TRUSTED

Seven checkpoints between an AI draft and a real-world decision.

The model prepares. The system checks. A person decides.

1. Signed-in identity (Entra)
2. Role and case scope
3. Permitted evidence only, with source and version IDs
4. Bounded model call, finite allowance and expiry
5. Schema, citation and freshness checks: stale or error stops
6. Human review with current permissions: accept or reject with a reason
7. Separate domain decision and audit trail (consent, quality, Finance)

Failure never becomes approval: the work falls back to its owner
Hazards bypass the AI path entirely
Unknown, stale, disabled or over-allowance → owned manual work · hazards go straight to people

Sources: controlled AI flow (diagram 04, E09) · ADR 0002 · AMA Plan #105, #106

S17 · Answers: B07 · B16 · B21

ACT 4 · BUILT TO BE TRUSTED

Model inference. Measure accepted value.

INFERENCE PLANNING SCENARIO

USD 0.44/home

USD 6,873 / 15,600 homes

20 tasks x 13,768 tokens
x USD 1.60/million

Accepted cost

Unknown

No accepted-outcome
unit cost established

Include review effort
and assurance

Double workload

USD 0.88/home

Inference only. Full project TCO and accepted-outcome cost remain unknown.

CODED

finite calls, tokens and expiry per authority
usage captured on the assistant path
Azure budget alerts

HELD

coded money budgets per task and period
per-request call and depth limits
accepted-outcome cost: not yet measured

E21: assumed workload; secondary price, 23 Oct 2026, not current verified price. E19: usage only. E05: controls coded; #435 not accepted.

S18 · Answers: B05 · B25 · B26 · B27

ACT 5 · THE PLAN

The delivery lesson is now being tested.

RUNNING PILOT | Registered OSF base xcpt4 | Unregistered exploratory amendments

Does consolidation preserve usable, changeable software?

Juan Manuel Servera Bondroit

Conception, methodology and evaluation.
Historical case analysis and manuscript.

Isabel Cabezas Martin

Independent execution / data collection.
Software Engineer, Frontier Company
at Microsoft. Co-author.

Separation is intended to reduce experimenter bias.
Both authors work at Microsoft; this is not independent institutional validation.

Multiple implementation epochs. Experimental outcome results are not reported here.

OSF base xcpt4: owner-confirmed 7 Nov. Current amendments unregistered; e71037d9, no tag. Detail: B13.

Future: connected history and better agent context.

FUTURE OPTIONS ONLY / Neither Fabric nor Power BI is a current Zava Forward component.

Portfolio history, reused

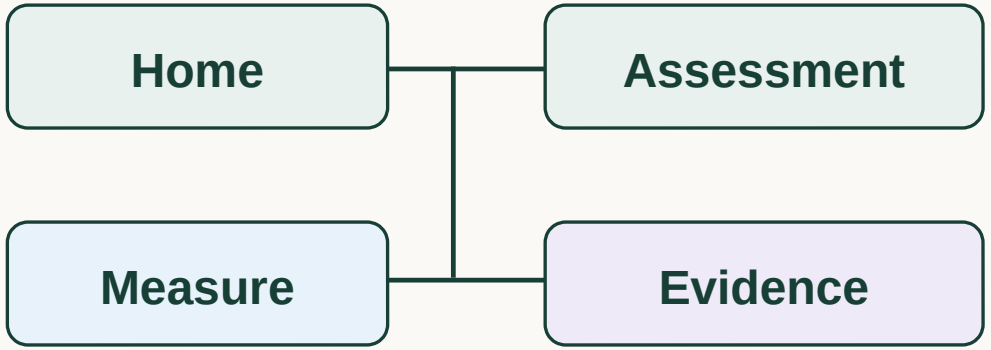
Join evidence across time.
Share reporting definitions and measures.

Curated facts -> analytics -> staff reports

Microsoft Fabric **Power BI**

Richer agent context

FABRIC ONTOLOGY / PREVIEW



60,000 fictional homes; not accounts

980 cases in the retained programme

43.8M synthetic readings / 730 days

WHY LATER | Prove the value, permissions, curated data and full cost first. Scale alone proves no platform need.

Sources: Fabric assessment + 30 Oct ontology follow-up. Preview; no integration or value gain proven.

Prove the need. Choose the smallest next step.

PRESERVE THE CURRENT FOUNDATION | Evaluate a need, not a service count.

1 Prove the current journey

Access, recovery and review
Resident and staff effort
Relevant actual evidence

2 Compare portfolio insight

SQL/read-model vs Fabric/PBI
Same cohort and history
Quality and complete cost

3 Test richer agent context

Current context vs ontology
Checked answers and refusals
Stale and forbidden-link tests

Agree owners. Cost the plan. Authorize separately. Expand only on evidence.

Sources: measurement design and acceptance boundaries. Separate execution authority required.

ACT 5 · THE PLAN

See it work: one home, one exception, one accountable next step.

1



Watch the evidence conflict find its owner

2



Watch what AI drafts, and what only a person can approve

3



Watch the record keep the next action



Ten-minute demo • route confirmed at preflight: live exercise or labelled evidence walkthrough

Sources: demo runbook (E07, E13)

S22 · Answers: B18

ZAVA FORWARD

More homes ready, with control intact.

1. Agree the validation gate
2. Name the owners
3. Review a costed plan



APPENDIX · NAVIGATION ONLY · END OF THE SPOKEN ROUTE

Questions we expect, and where the answers are.

STRATEGY AND VALUE

B01 · 29 Why not fix the process and hire more people?
B08 · 36 Will this double delivery and capture GBP 40m?
B09 · 37 What is the per-home case and cash exposure?
B17 · 45 When can we go live, and who owns next?
B20 · 48 When would Fabric add value beyond SQL?

METHOD AND READINESS

B12 · 40 Validated with real residents and staff?
B13 · 41 Who runs the study, under which protocol?
B15 · 43 What must happen before a pilot decision?

COST AND TOKENOMICS

B04 · 32 Where do rules, models and generation fit?
B05 · 33 What does the technology cost today?
B25 · 53 Tokenomics unit: can the AI bill run away?
B26 · 54 Why not cache everything, use the cheapest model?
B27 · 55 What stops chat and agent sprawl?

ARCHITECTURE AND AI

B02 · 30 Replacing or integrating our housing system?
B03 · 31 Which autonomy levels did we meet or change?
B10 · 38 How is the application architecture organised?
B11 · 39 What proves the AI is real and useful?
B14 · 42 Is it agentic if so much is deterministic?

TRUST, RISK AND PEOPLE

B06 · 34 How do residents keep choice, incl. offline?
B07 · 35 What stops humans rubber-stamping AI?
B21 · 49 How do you prevent hallucination and injection?
B22 · 50 Can data leave the UK or train a model?
B23 · 51 How will you know it is fair and accessible?
B24 · 52 How do policy, law and funding stay current?

OPERATIONS AND DEMO

B16 · 44 What if the model, cloud or deployment fails?
B18 · 46 Can you demo a new live action now?
B19 · 47 How do device readings support safe decisions?

Why not fix the process and hire more people instead?

OPTION 1

Manual / process only

Process improvement is the baseline

TRADE-OFF

Additional staff or contractor capacity may still be necessary

STAFFING PLAN SET IN PILOT

OPTION 2

Workflow + rules

Workflow must beat the baseline

TRADE-OFF

If rules and templates deliver the better accepted-task result, retain them

COMPARISON 1 · READY TO RUN

OPTION 3

Workflow + bounded AI

AI must beat the same workflow without AI

TRADE-OFF

If rules and templates win, report the negative AI finding

COMPARISON 2 · READY TO RUN

ALL OPTIONS

Physical delivery capacity still limits throughput

Sources: E02 · E10

Is this replacing our housing system or integrating with it?

CURRENT DEMONSTRATOR

Controlled workbench

Demonstrator of coordination and evidence contracts

Not a replacement decision

Housing and repairs integration is the next step

INTEGRATION ROADMAP

- **Housing system**
- **Repairs system**
- **Configured or designed M365 path**

Integration scope, authority and migration need separate discovery

PRODUCTION APPRAISAL MUST IDENTIFY

Authoritative systems · field ownership · APIs · event and reconciliation contracts
Security · retention · the manual service route

Sources: E06 · E07

Which autonomy levels did we meet or change?

BUSINESS AGENT	BRIEF	CURRENT	WHY / REMAINING WORK
Stock Prioritisation	L2 Recommend	L2 recommend Boundary matches	Rules + stock facts; full EPC/health/grant cohort policy remains partial
Assessment	L3 draft Coordinator approves	L2 preparation Scope reduced	Survey/design checklist, not a full PAS assessment; professional review retained
Grant & Funding	L3 draft Finance submits	L2 preparation Scope reduced	Application/claim gaps, not a full bid; Finance still owns submission
Works + residents Coordination	L3 acts Within preferences	L2 options Authority reduced	No auto-book/send. Prove capacity, preferences and current confirmation
Quality Verification	L4 autonomous Flags only	L2 concerns Authority reduced	No autonomous QV domain flag; independent QA and OD-31 trial remain

L3 draft/review rights are preserved; full drafts are not delivered. Automatic invocation != decision authority.

Sources: original brief; current assistants/ranking; PRD FR-047; BRD OD-31/32. Reviewed 8 Nov 2026.

Where do rules, decision models and generation fit?

01

Rules / template

Exact labelled fields
Arithmetic and ranking
Deterministic gates

THE SIMPLER BASELINE

02

Decision 1

Supplied candidates
Binary support / abstain
#58 only; rebuild pending

MUST BEAT A PARSER

03

Generative model

Sourced narrative
Permitted evidence only
Human checks meaning

NOT AN APPROVAL

Same cases. Held-out quality. Total human effort. Latency. Full cost.

A successful call proves integration, not lower cost, faster work or calibrated accuracy.

Sources: E04 · E05 · E15

What does the technology cost today?

USD 700 owner-set October all-in forecast ceiling

Execution gates remain; authority expires 25 October 16:02 UTC

Refresh post-change forecast; outside-group cost remains unknown

Amount	Scope	Class	Status
USD 700	October all-in forecast ceiling, owner decision 22 Oct	Owner-set ceiling	Paid tranches blocked; expires 25 Oct 16:02 UTC
USD 242.36	Accrued 11-22 Oct, resource group ama	Observed, partial	Estimated until invoiced
USD 535.33	Pre-change scenario	Historical forecast	Superseded for authorisation; refresh
USD 450/month	EXA-012, 10 Oct	Historical allocation	Not current authority
Outside-group spend	Outside resource group	Unknown	Never zero by assumption

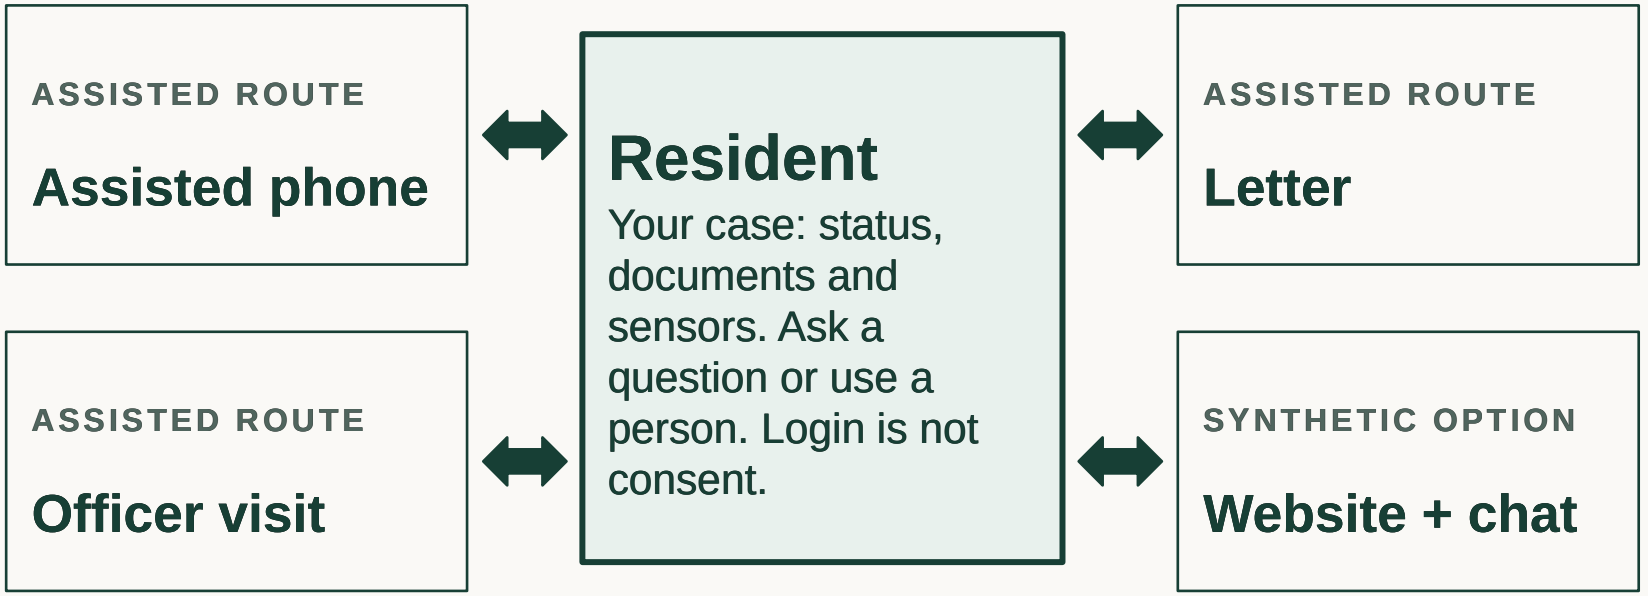
DATED LEDGER
NOT PRODUCTION TCO

Future project appraisal
P + USD 6,872.9856
Per planned home:
(P + USD 6,872.9856)
/ 15,600
P = non-inference costs for the same period.
P is not yet priced.

Count each cost once. Do not add the October USD 700 envelope and its included inference.

Sources: E11 / E05 / E21. Historical price assumption: 23 Oct 2026; no verified current quote.

How do residents retain choice, especially without digital access?



Optional website + chat; assisted routes stay

Offer ≠ agreement; silence ≠ consent

Representative authority is task-limited

SOURCE BUILT | AI OFF | HUMAN VALIDATION NEXT

AI tariff expired 1 Nov. Synthetic callbacks are not real bookings. Resident validation remains.

Sources: E02 · E08 · E12

What stops a human from rubber-stamping AI?

Source verification before proposal acceptance

Current permission, evidence and independent authority

Meaningful oversight remains a human acceptance gate

1. Capture

Role + case scope
Permitted evidence
Source/version IDs

AI PREPARES

2. Call

Bounded task
Model + usage
No database lock

3. Validate

Schema + citations
Stale/error = stop
Store as unreviewed

PEOPLE DECIDE

4. Review

Human checks source
Current permissions

Separate domain action

Failure does not become approval

Unknown, stale, disabled, exhausted allowance or unavailable model → owned manual work; hazards bypass AI.

Citation presence is not semantic truth · Proposal review is not consent, QA or Finance

SOURCE VIEW · 23 OCT 2026

Sources: E02 · E09 · E12

Will this double retrofit delivery and capture GBP 40 million?

1 · AMBITION

Original brief ambitions: the outcomes the pilot measures

Double retrofit delivery

Capture GBP 40 million

ORIGINAL BRIEF AMBITION

2 · DEPENDENCIES

Capacity, scope, eligibility and cash all constrain value

Survey, technical, contractor, material, resident and QA capacity

Grant availability and secured cash: confirmed at the funding gate

3 · TEST

Frozen cohorts and attributable comparisons required

Measure accepted handovers and the whole cost of the work

READY TO RUN

4 · OBSERVED RESULT

NEXT MILESTONE

Benefit range set once the baseline, causal limits and sensitivity are agreed

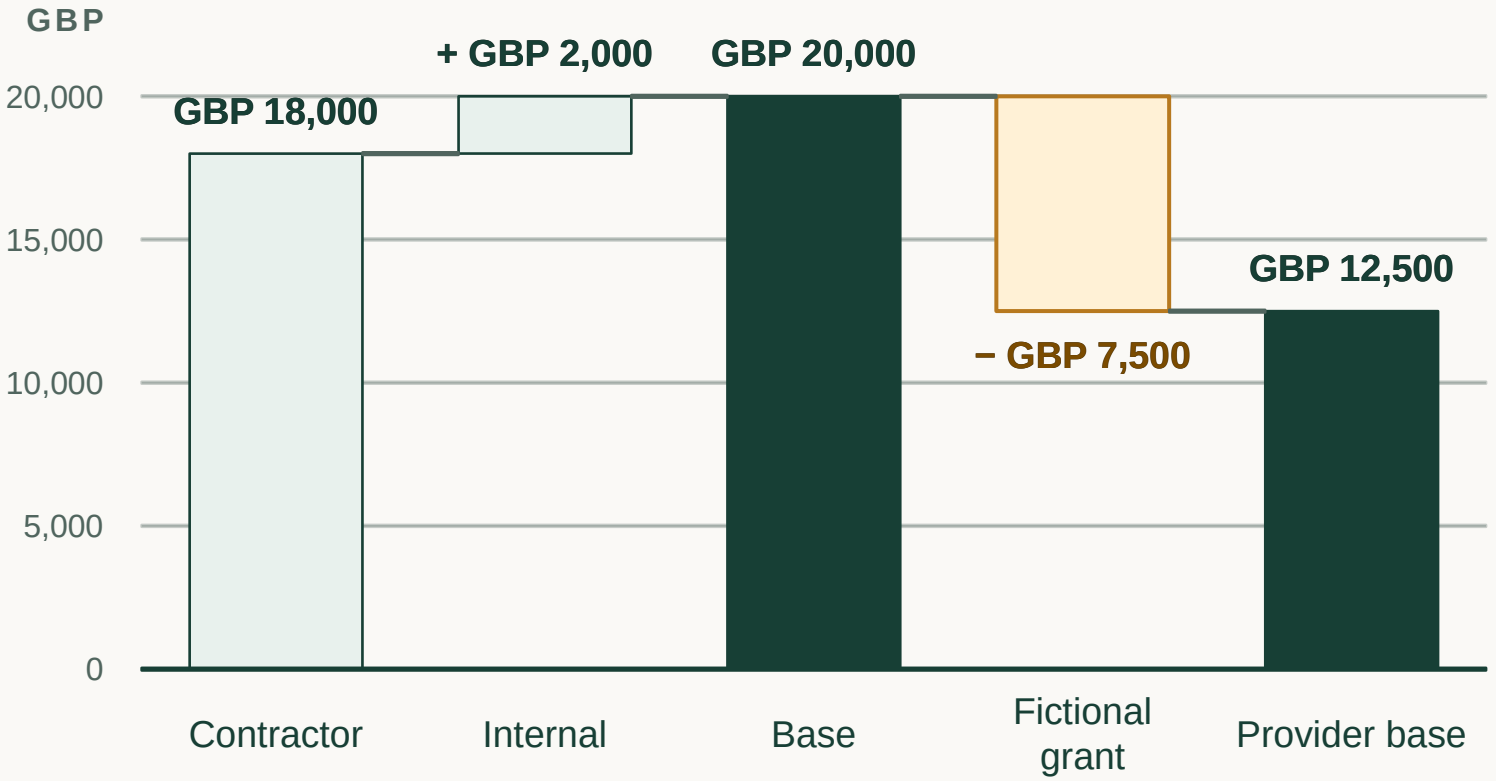
Saved minutes become cash through a Finance-approved mechanism

Sources: E01 · E02 · E10

What is the per-home business case and cash exposure?

Benchmark: GBP 6,082/home
x 15,600 = GBP 94,879,200 works

Conditional: min(50% x GBP 6,082, GBP 7,500)
= GBP 3,041/home; x 15,600 = GBP 47,439,600



FICTIONAL EXERCISE ARITHMETIC

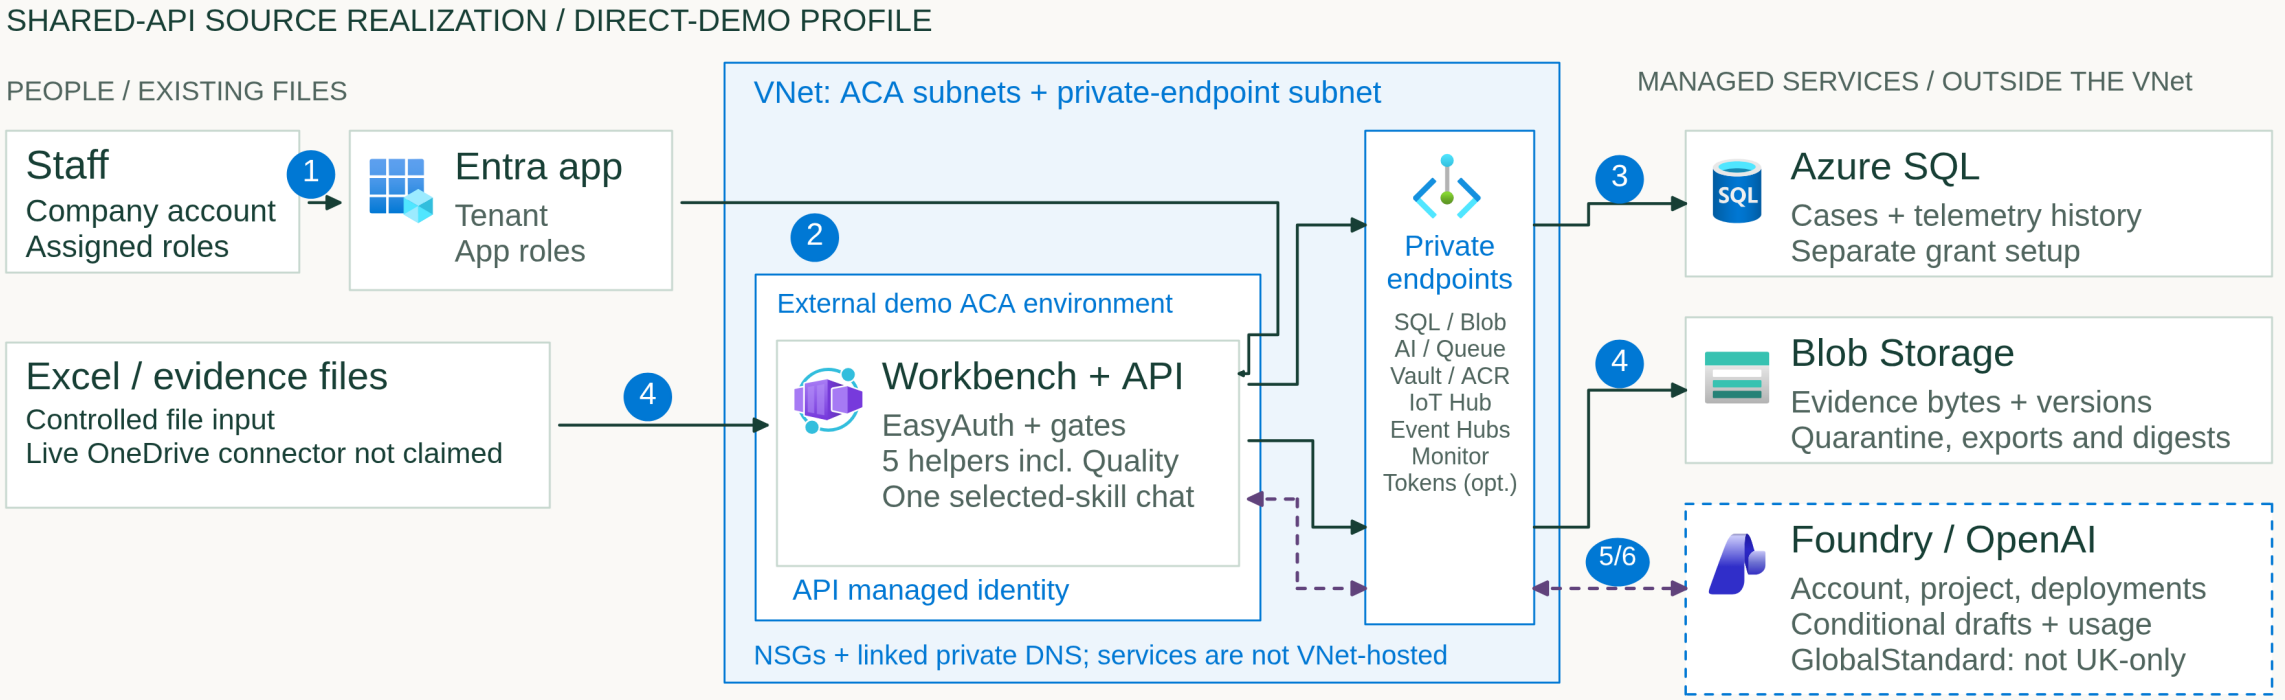
Fictional base GBP 20,000
GBP 22,000 with contingency

Gross cash GBP 20,000 before grant
Fictional receivable GBP 7,500
An award is not cash.

Contingency spent: provider GBP 14,500
Evidence supports claims; not eligibility, award or payment.

E02 / E21: historical benchmark, not quote. Grant assumes eligible homes/costs, allocation, award and claim; Finance/funder review required.

How is the full application architecture organised?



Detail: full-size editable map. Sensors / averages / resident / support: main slide 13. People own consent, sign-off and release.

Core view here. Full editable reference: device flow, operations, production/worker options and separate resident boundary.

Sources: Bicep + current source contracts, 8 Nov 2026. Gated paths are not live acceptance.

What proves the AI is real and useful rather than a template?

DATED HISTORICAL EVIDENCE

Invocation
Provider/model/usage and source=ai evidence
Historical real calls and persistence are documented.

≠

NEXT MILESTONE

Semantic acceptance
Semantic quality, critical omissions, unsafe advice and fairness require separate evaluation.

NEXT MILESTONE

Incremental value
Review burden and incremental value require separate evaluation.

LABELLED FALLBACK

Template/replay is clearly labelled as a fallback
The real-inference gate runs on live model calls.

WHAT THE PILOT ADDS

Real invocation proven; calibrated quality and value come from the pilot
Live readiness is re-checked before every run, never assumed.

Sources: E07 · E09 · E04 · E13

Did you validate this with real residents and staff?

Not yet. This is desk-research synthesis, not resident endorsement.

ZAVA FORWARD / PROPOSED RESPONSES

PRODUCTS / SERVICES

Owned case
Visit + help

GAIN CREATORS

Scope-specific choice
Agreed arrangements

PAIN RELIEVERS

One named contact
Visible open remedies

FIT TO TEST

Not proven benefit

RESIDENT / HYPOTHESISED PROFILE

GAINS SOUGHT

Meaningful choice
Owned aftercare

JOBS

Choose works
Keep routines
Use / query

PAINS TO TEST

Repeating a story
Pressure to agree

Next: direct task and comprehension tests, assisted routes, real handoffs and total effort.

Sources: resident journey atlas, 14 Oct / reconciliation 23 Oct; source-linked DT trace, 30 Oct. Hypotheses only.

What is running, and how independent is the execution?

3 models x 5 arms x 21 = 315 planned trajectories; F1-F10 + held-out CR1.

GPT 6.1 Sol / Claude Sonnet 5.5 / Gemini 4 Argon; medium effort; seed 2027. Active 21, default 12.

ARM / TREATMENT	
A	Feature-only delivery
B	Scheduled consolidation
C	Read-only review
D	Matched-resource longer delivery
E	Visible regression repair

IMPLEMENTATION / COMPLETED STEPS	
99a1ed1	651
3c2ce89	227
e71037d9 (no tag)	369

1,247 steps; owner-reported progress

Not independent n; not one frozen implementation.

Execution separated from the designer; not independent analysis or institutional validation. No outcome results reported.

MVE source e71037d9; owner-confirmed 8 Nov profile. OSF base xcpt4; current amendments unregistered.

Is this actually agentic if so much is deterministic?

Controlled workflow

RULES · NO AI

Rules rank, route and protect state.
A model cannot advance a case.

Automatic preparation

APPROVED SCANS / L2

Gap scan -> proposal -> human review.
One model call per attempt; no tools.

Accountable decision

PEOPLE DECIDE

Resident, professional and Finance rights stay.
Chat answers; decisions remain separate.

AUTONOMY

Automatic invocation != decision authority.
L3 actions and L4 flags need separate evaluated activation.

Sources: E03 · E06 · E09

What exactly must happen before a pilot or production decision?

NEXT GATE

Current deployment and AI evidence

Current deployment, roles, SQL and real-AI evidence

OWNER ROLE

Technical owner

REQUIRED ARTIFACT

Staff access and role denial, task-bound real inference, SQL persistence on the selected deployment

NEXT GATE

Human and control acceptance

Human, privacy, professional and accessibility acceptance

OWNER ROLE

Professional, privacy, accessibility and safety reviewers

REQUIRED ARTIFACT

Task/comprehension, accessibility, safety/privacy and professional acceptance

NEXT GATE

Value, recovery and release

Measured value, cost, recovery and explicit release decision

OWNER ROLE

Programme and Finance

REQUIRED ARTIFACT

Measured value and cost against the comparator, recovery test, explicit release decision

Sources: E02 · E12 · E13 · E17

What happens if the model, cloud or deployment fails?

Resilience is configured in the architecture; whole-system recovery drills are the next milestone.

Access + secrets

Entra and managed identities
Key Vault secret references
Private backing services

Runtime + releases

Container Apps / pinned images
Registry and CI checks
Controlled rollout + rollback

Observe + budget

Monitor / Log Analytics
Scoped allowance + usage
Budget alerts warn; hard stops are designed.

Evidence required before expansion

NEXT GATE

Genuine staff sign-in + role denial	Restore and recovery drill
Azure SQL persistence	Human task/comprehension checks
Bounded real-model result	Current cost + processing terms

Restart proven; full restore drill is the next milestone

OWNED MANUAL ROUTE

No approval from failure or stale evidence
Owned manual work and independent safety handling
Keep manual and safety routes independent; do not relax private controls to rescue a demo.

Columns are concerns, not request hops · Amber = next milestone

SOURCE VIEW · 23 OCT 2026

Sources: E06 · E07 · E13

When can we go live and who owns the next step?

Dependency-led plan: dates follow the gates

SEQUENCED PLAN

READINESS SEQUENCE	OWNER	
1 Agreed validation scope	Programme	Programme owns scope; Finance owns spend review
2 Refreshed cost	Finance	
3 Deployment prerequisites	Technical owner	Technical and control owners supply acceptance evidence
4 Finite test authorised	Programme + Finance	
5 Dates set	→	LATER, SEPARATE DECISIONS: Commercial pilot; Operational release

Can you demonstrate a new live action right now?

DEMO BOUNDARY

Only pre-authorised, reversible constructed scope
No unplanned role, network, spend or data changes
We will not grant a role, bypass private networking, enable assistant authority, submit a grant or send a resident message on the fly.

DEMO ROUTE, CONFIRMED AT PREFLIGHT

→ **L · Live, preflight passed**
Only within the preflighted, authorised demonstration boundary

→ **E · Repository evidence**
Fallback is visibly labelled evidence, never simulated live proof
Switch once; say exactly what it does and does not prove.

→ **U · Unavailable**
Explain the control instead

□ **No silent mock fallback**

Sources: E07 · E13

How do device readings support decisions safely?

SOURCE-ONLY TARGET FLOW / LIVE END-TO-END HUB + SQL ACCEPTANCE UNPROVED

01	02	03	04	05
Synthetic device	**IoT Hub Basic**	**Private Event Hub**	**Intake + SQL**	**Human review**
Exact binding Stable reading ID	Device boundary One-way intake	Hub-only sender Keyless receiver	Accepted readings On-read averages	Review the signal No diagnosis

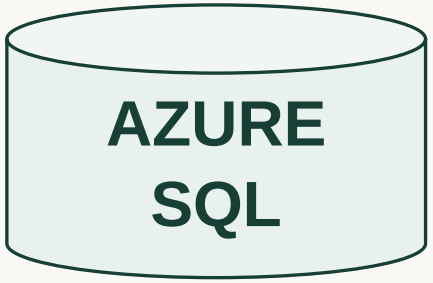
Why Basic: device identity + one-way intake. No Standard device-control features needed.

Hazard reports go straight to human safety. Neither telemetry nor AI may delay the response.

Sources: E12

When would Fabric add value beyond SQL?

RETAIN / OPERATIONAL CONTROL



Transactions
Ledger and audit
Accountable workflow

Keep Excel / M365 as bounded inputs.
Compare full integration and control costs.
Fabric SQL lacks ledger creation;
a warehouse is not the write workflow.

FUTURE CHOICE / SELECTIVE ANALYTICS

Reusable cross-home history.
Shared measures and portfolio reports.

Curated facts → Analytics → Staff reports

Compare with SQL views + Power BI.
Require better effort, freshness or total cost.

NO PURCHASE OR EXPORT APPROVED | Access, history and total cost need separate review.

Sources: E02 · E03 · E06

How do you prevent hallucinations and prompt injection?

Permitted task context and source/version binding

Schema/citation checks plus independent semantic review

Built to contain vulnerabilities and hallucinations, not to deny them

1. Capture

Role + case scope
Permitted evidence
Source/version IDs

2. Call

Bounded task
Model + usage
No database lock

3. Validate

Schema + citations
Stale/error = stop
Store as unreviewed

4. Review

Human checks source
Current permissions
Separate domain action

UNTRUSTED EVIDENCE

Documents are evidence, not instructions

REVIEW BOUNDARY

Citation presence is not semantic truth

Failure does not become approval

Unknown, stale, disabled, exhausted allowance or unavailable model → owned manual work; hazards bypass AI.

Filled arrow = guarded request/response step, not agent-to-agent talk

Sources: E09 · E12 · E07

SOURCE VIEW · 23 OCT 2026

Can customer data leave the UK or train a provider's model?

Processing geography is set by the deployment type

UK-only inference is a deployment choice (a regional deployment rather than GlobalStandard), confirmed at the security gate.

Verify exact model, deployment type and contract

before real data is processed

- VERIFY Provider
- VERIFY Model and version
- VERIFY Deployment SKU
- VERIFY Routing and geography
- VERIFY Retention
- VERIFY Abuse monitoring
- VERIFY Subprocessors
- VERIFY Applicable agreement

No real resident data or private dataset in this handoff

The handoff contains only public synthetic illustrations and repository-authored summaries, not private datasets or credentials.

Sources: E04 · E06 · E12 · E13

How will you know this is fair and accessible?

No-digital route and understandable explanations

Cohort errors, exclusions and unknowns are visible

Engineering checks ≠ resident or WCAG acceptance

	Task success	Comprehension	Errors and exclusions by cohort
Digital route	MEASURED IN PILOT	MEASURED IN PILOT	MEASURED IN PILOT
No-digital route	MEASURED IN PILOT	MEASURED IN PILOT	MEASURED IN PILOT

Inspect errors and exclusions by agreed cohorts without exposing small groups

GUARDRAIL

An AI result that improves average speed while worsening a critical group outcome should not pass the gate.

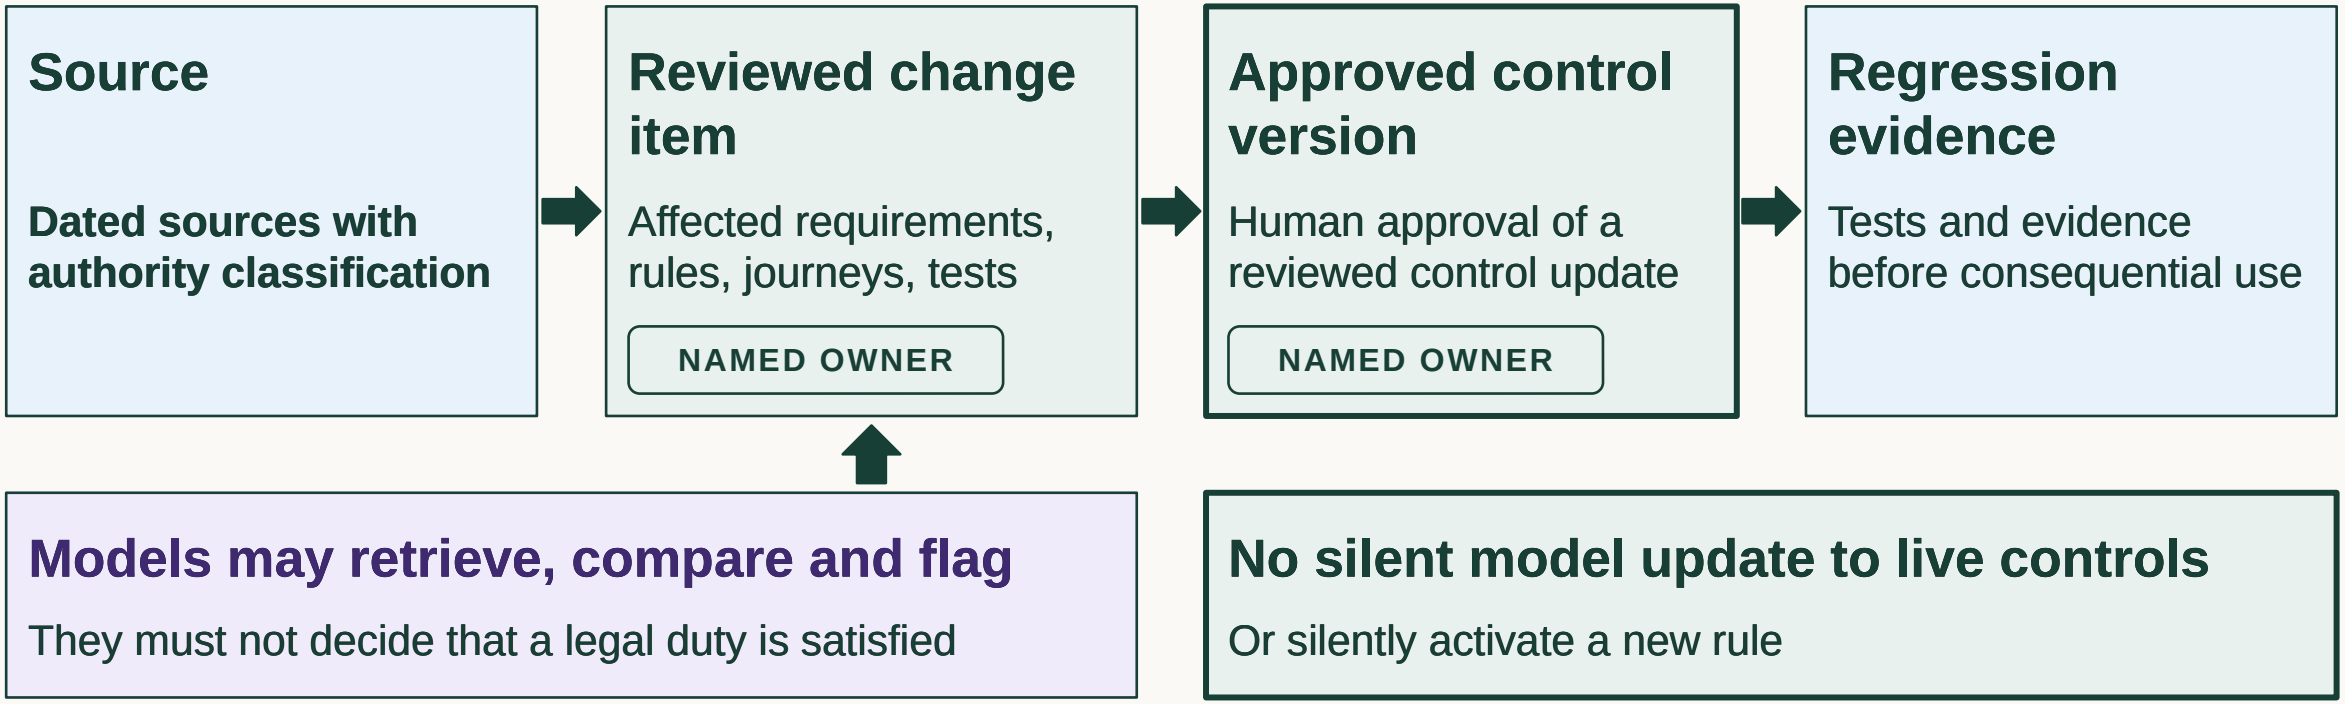
Conformance and fairness are assessed with real users and a full audit in the pilot.

Sources: E02 · E12 · E15

How will you keep policy, law and funding rules current?

Detected change → reviewed rule change → tests

REPOSITORY CONVENTION · READY TO ACTIVATE



Dated official-source snapshots; currentness and named legal review required before consequential use.

Sources: E02 · E12 · E18

What is your Tokenomics unit and can the AI bill run away?

Planning: USD 0.440576/home inference; accepted-outcome TCA unknown

Historical secondary price: 23 Oct 2026. Assumed workload, not current verified deployment price.

Observe, then assume

6,884 / 5 = 1,376.8 tokens/run
x 10 = 13,768 tokens/task
x 20 = 275,360 tokens/home

Price inference only

275,360 / 1m x USD 1.60
= USD 0.440576/home
x 15,600 = USD 6,872.9856

Measure accepted TCA

Hosting + review + assurance
Accepted above quality floor
TCA / accepted work: unknown
Double workload: USD 0.881152/home

Inference scenario is not total project cost or accepted value

SOURCE-CODED · GATED

Existing bounded path

Finite calls / tokens / expiry; partial metering
Coded money / request-depth admission.

Live gates: not run

Full attribution; live money-stop proof
Request-bound proof; billing reconciliation.

Tokenomics framing: our working method, refined with Finance

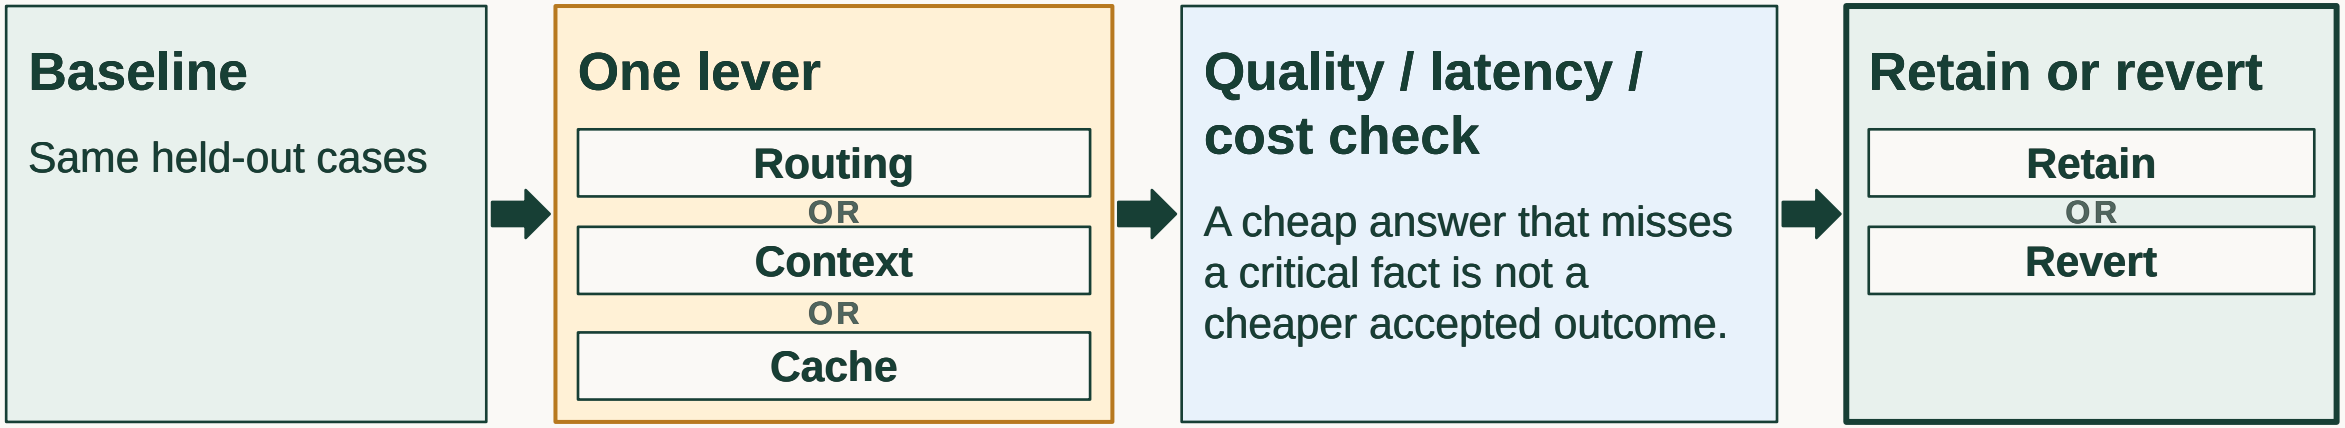
RECOVERY: 24 OCT 2026

Sources E19 / E05 / E11 / E21. #435 not accepted; no saving, payback or full-project price established.

Why not cache every answer and route everything to the cheapest model?

Change one lever, compare, then keep or revert

OPTIMISATIONS · ON THE ROADMAP



CACHE

Same permission, evidence version and task validity required

ROUTING

Smallest model must clear held-out quality and latency thresholds

Sources: E19 · E04 · E09

What stops chat and agent sprawl multiplying hidden calls?

USER QUESTION

One selected skill. One bounded call.

No routing-model hop, tool loop or write tool.

BEFORE DISPATCH

Scope and release Financial admission

Unknown authority or liability: stop; keep manual work.

ADVISORY OUTPUT

Advice, not a case decision.

Sources and permissions checked again.

SOURCE-CODED BOUNDS

Calls/tokens/expiry + task/period money limits.
Durable request call/depth ceilings.
Approved evidence required before dispatch.

LIVE ACCEPTANCE STILL SEPARATE

Billing and whole-journey coverage unproved.
Historical calls and resident accounting remain explicitly unjoined.

Sources: API bounded-skills and core consumption guides, 8 Nov 2026. Source controls, not live financial acceptance.